



Cairo University
Faculty of Engineering
Computer Engineering Department



OS Scheduler

Document Structure

- Objectives
- Introduction
- System Description
- Guidelines
- Grading Criteria
- Deliverables

Objectives

- Evaluating different scheduling algorithms.
- Practice the use of IPC techniques.
- Best usage of algorithms, and data structures.

Platform *Linux*

Language *C*

Introduction

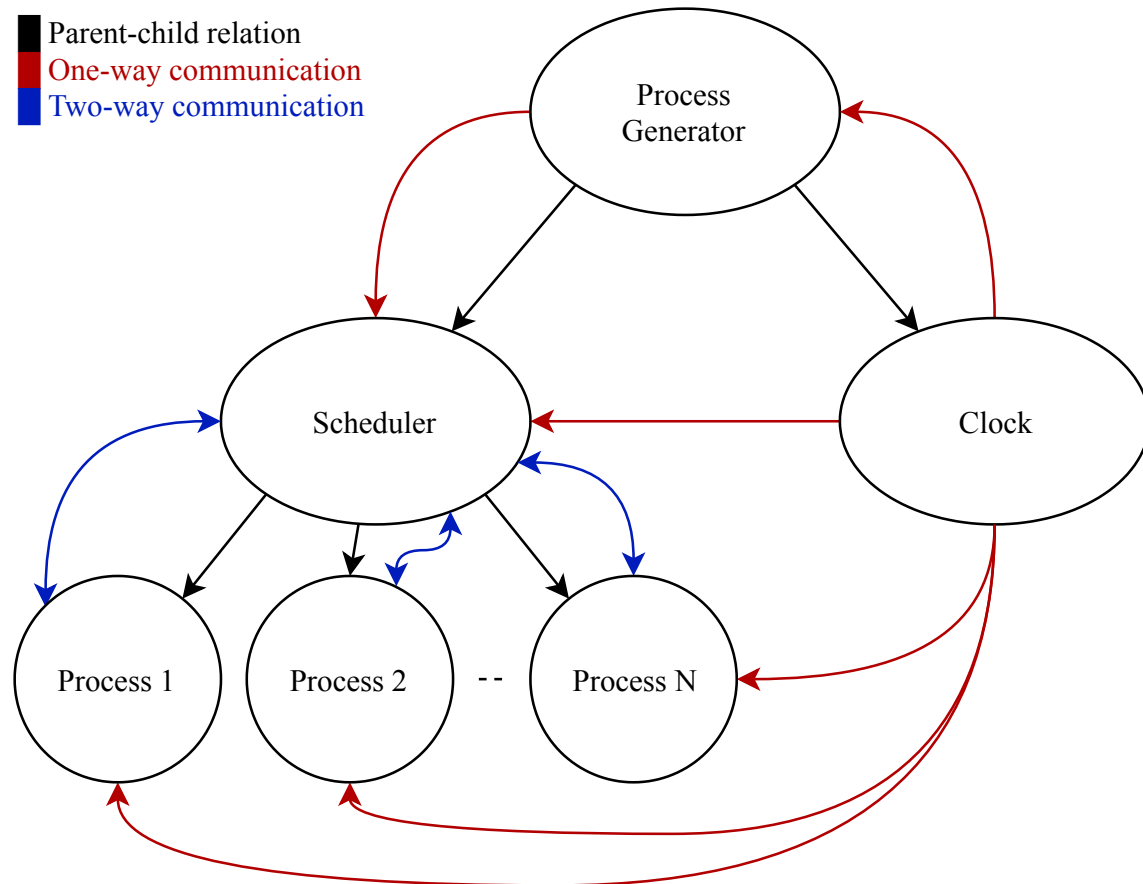
A CPU scheduler determines an order for the execution of its scheduled processes; it decides which process will run according to a certain *data structure* that keeps track of the processes in the system and their status.

A process, upon creation, has one of the three states: *Running*, *Ready*, *Blocked* (doing I/O, using other resources than CPU or waiting on unavailable resource).

A bad scheduler will make a very bad operating system, so your scheduler should be as much optimized as possible in terms of memory and time usage.

System Description

Consider a Computer with 1-CPU and infinite memory. It is required to make a scheduler with its complementary components as sketched in the following diagrams.



Part I: Process Generator (Simulation & IPC)

CODE FILE *process_generator.c*

The process generator should do the following tasks...

- Read the input files (check the input/output section below).
- Ask the user for the chosen scheduling algorithm and its parameters, if there are any.
- Initiate and create the scheduler and clock processes.
- Create a data structure for processes and provide it with its parameters.
- Send the information to the scheduler *at the appropriate time* (when a process arrives), so that it will be put in its turn.
- At the end, clear IPC resources.

Part II: Clock (Simulation & IPC)

CODE FILE *clk.c*

The clock module is used to emulate an integer time clock. *This module is already built for you.*

Part III: Scheduler (OS Design & IPC)

CODE FILE *scheduler.c*

The scheduler is the core of your work, it should keep track of the processes and their states and it decides - based on the used algorithm - which process will run and for how long.

You are required to implement the following THREE algorithms...

1. Non-preemptive Highest Priority First (HPF).
2. Shortest Remaining time Next (SRTN).
3. Round Robin (RR).

The scheduling algorithm only works on the processes in the *ready queue*. (Processes that have already arrived.)

The scheduler should be able to

1. Start a new process. (Fork it and give it its parameters.)
2. Switch between two processes according to the scheduling algorithm. (Stop the old process and save its state and start/resume another one.)
3. Keep a *process control block (PCB)* for each process in the system. A PCB should keep track of the state of a process; running/waiting, execution time, remaining time, waiting time, etc.
4. Delete the data of a process when it gets notified that it finished. *When a process finishes it should notify the scheduler on termination, the scheduler does NOT terminate the process.*
5. Report the following information
 - (a) CPU utilization.
 - (b) Average weighted turnaround time.
 - (c) Average waiting time.
 - (d) Standard deviation for average weighted turnaround time.
6. Generate two files: (check the input/output section below)
 - (a) Scheduler.log
 - (b) Scheduler.perf

Part IV: Process (Simulation & IPC)

CODE FILE *process.c*

Each process should act as if it is CPU-bound.

Again, *when a process finishes it should notify the scheduler on termination, the scheduler does NOT terminate the process.*

Part V: Input/Output (Simulation & OS Design Evaluation)

Input File

| <i>processes.txt</i> example | | | |
|------------------------------|---------|---------|----------|
| #id | arrival | runtime | priority |
| 1 | 1 | 6 | 5 |
| 2 | 3 | 3 | 3 |

- Comments are added as lines beginning with `#` and should be ignored.
- Each non-comment line represents a process.
- Fields are separated with *one tab character* `'\t'`.
- You can assume that processes are sorted by their arrival time. *Take care that 2 or more processes may arrive at the same time.*
- You can use the *test_generator.c* to generate a random test case.

Output Files

| <i>scheduler.log</i> example | | | | | | | | | | | |
|------------------------------|---------|----|---------|---|----------|-----|---|-------|---|--------|-------------------------|
| # | At time | x | process | y | state | arr | w | total | z | remain | y wait k |
| | At time | 1 | process | 1 | started | arr | 1 | total | 6 | remain | 6 wait 0 |
| | At time | 3 | process | 1 | stopped | arr | 1 | total | 6 | remain | 4 wait 0 |
| | At time | 3 | process | 2 | started | arr | 3 | total | 3 | remain | 3 wait 0 |
| | At time | 6 | process | 2 | finished | arr | 3 | total | 3 | remain | 0 wait 0 TA 3 WTA 1 |
| | At time | 6 | process | 1 | resumed | arr | 1 | total | 6 | remain | 4 wait 3 |
| | At time | 10 | process | 1 | finished | arr | 1 | total | 6 | remain | 0 wait 3 TA 10 WTA 1.67 |

- Comments are added as lines beginning with `#` and should be ignored.
- Approximate numbers to the nearest 2 decimal places, e.g. 1.666667 becomes 1.67 and 1.3333334 becomes 1.33.
- Allowed states: *started*, *resumed*, *stopped*, *finished*.
- TA & WTA are written only at *finished* state.
- You need to stick to the given format because files are compared automatically.

| <i>scheduler.perf</i> example | |
|-------------------------------|--------------------|
| CPU | utilization = 100% |
| Avg | WTA = 1.34 |
| Avg | Waiting = 1.5 |
| Std | WTA = 0.34 |

- If your algorithm does a lot of processing, processes might not start and stop at the same time instance. Then, your utilization should be less than 100%.

Guidelines

- Read the document carefully at least once.
- You can specify any other additional input to algorithms or any assumption but after taking permission from your TA.
- The user should be able to choose between different scheduling algorithms.
- You should specify how your algorithm handles ties.
- Priority values range from 0 to 10 where *0 is the highest priority and 10 is the lowest priority*.
- Your program must not crash.
- You need to release all the IPC resources upon exit.
- The measuring unit of time is 1 sec, there are no fractions, so no process will run for 1.5 second or 2.3 seconds. Only integer values are allowed.
- You can use any IDE (Eclipse, Code::Blocks, NetBeans, KDevelop, CodeLite, etc.) you want of course, though it would be a good experience to use make files and standalone compilers and debuggers if you have time for that.
- Spend a good time in design and it will make your life much easier in implementation.
- The code should be clearly commented and the variables names should be indicative.

Grading Criteria

- NON compiling code = ZERO grade.
- Correctness & understanding (50%).
- Modularity, naming convention, code style (20%).
- Design complexity & data structures used (20%).
- Team work (10%).

Deliverables

You should deliver code files, test cases and report containing the following information...

- Data Structure used.
- Your algorithm explanation and results.
- Your assumptions.
- Workload distribution.
- A table for time taken for each task. *It will not affect your grade so please be honest.*

Keep the document as simple as possible and do not include unnecessary information we do not evaluate by word count!